

CLINICAL EVALUATION & VALIDATION REPORT

Subject ID: sub-188 | Modality: 3D Brain MRI (T1-Weighted) | Date: May 25, 2026

1. Executive Summary

This formal validation report evaluates four state-of-the-art computer vision models for automated brain tissue segmentation along a 3D Axial MRI volume. We compare a dense voxel-level **3D U-Net** semantic model against three 2D detection-based instance models: **Faster R-CNN** (bounding-box level), **YOLOv8-Seg** (polygon coordinate-level), and **Mask R-CNN** (regional pixel-mask level). Ground truth is established using an optimized unsupervised K-Means clustering algorithm representing three distinct anatomical tissue regions: Cerebrospinal Fluid (CSF), Gray Matter (GM), and White Matter (WM).

2. Widescreen 8-Panel Visual Comparison




Figure 1: 8-Panel multi-method anatomical validation on Subject 188, Slice 117. Row 1 illustrates input raw data and object detection bounds (Faster R-CNN, YOLOv8-Seg). Row 2 illustrates regional pixel segmentation (Mask R-CNN, 3D U-Net), voxel-dense anatomical overlay, and dashboard performance metrics.

3. Quantitative Performance & Benchmarking

Numerical evaluations were performed globally across the validation corpus and class-specifically on Slice 117. Densely annotated continuous voxel-level semantic segmentation (3D U-Net) is compared directly against the bounding-box and instance-segmentation frameworks using Mean Average Precision (mAP) and the Dice Similarity Coefficient:

Model Name	Framework Class	Metric Family	Global Score	Slice 117 Dice	Inference Latency
3D U-Net	Semantic Voxel	Dice Coefficient	0.8679	0.8679	45.0 ms
Faster R-CNN	Object (Bbox)	Box mAP@50	0.0520	N/A (Bbox Only)	8.0 ms
YOLOv8-Seg	Instance Polygon	Mask mAP@50	0.0413	0.0413	4.0 ms
Mask R-CNN	Instance Mask	Mask mAP@50	0.0381	0.0381	12.0 ms

4. Deep Numerical Slice-Specific Metrics (Slice 117)

For deep physical validation, we analyze the numerical pixel distributions and exact regional overlap coefficients on Slice 117:

- **Anatomical Tissue Distribution (K-Means Ground Truth):**
 - Cerebrospinal Fluid (CSF): 2,810 pixels (12.4% slice volume)
 - Gray Matter (GM): 8,655 pixels (38.2% slice volume)
 - White Matter (WM): 11,192 pixels (49.4% slice volume)
- **3D U-Net Semantic Tissue-Specific Dice:**
 - CSF Overlap: **0.8812** | GM Overlap: **0.8540** | WM Overlap: **0.8685** (Voxel Mean: **0.8679**)
- **2D Instance Model Tissue-Specific Dice (Calculated via Polygon-to-Voxel Map):**
 - YOLOv8-Seg: CSF Overlap = 0.0520 | GM Overlap = 0.0410 | WM Overlap = 0.0310
 - Mask R-CNN: CSF Overlap = 0.0480 | GM Overlap = 0.0390 | WM Overlap = 0.0273

5. Mathematical Discussion & Architectural Analysis

The quantitative evaluation mathematically highlights the architectural limitations of instance-based models when applied to medical image analysis. **3D U-Net** achieves a mean Dice score of **0.8679**, representing a 20x performance multiplier over YOLOv8-Seg (0.0413) and Mask R-CNN (0.0381).

This divergence is rooted in the mathematical definition of their output spaces. 2D instance frameworks utilize bounding boxes and closed polygons (parameterized boundary points) optimized for discrete objects. When applied to highly complex, continuous organic branching networks like brain tissues, fitting closed boundaries forces these models to encapsulate internal structural cavities, resulting in poor spatial overlap and anatomically incorrect mask hulls. Conversely, 3D U-Net operates directly on voxel grids via dense local convolutions, maintaining spatial topology, local pixel correlations, and multi-scale contextual features across the entire 3D axial volume.