A Deep Learning Comparative Framework for Volumetric Tissue Segmentation and Diagnostic Classification of Brain MRI

cook45[1], and **clack**[2]

[1]*Senior Deep Learning Systems Architect, Department of Neuroimaging and Artificial Intelligence*

[2]*Lead Medical Imaging Systems Engineer, Division of Computer Vision & Diagnostic Systems*

Abstract

This article presents a comprehensive comparative analysis of deep learning architectures applied to clinical brain Magnetic Resonance Imaging (MRI). We address the two foundational tasks of automated neuroimaging analysis: diagnostic classification and high-resolution spatial tissue segmentation. For classification, we evaluate a Vision Transformer (ViT-Base) against a standard convolutional ResNet-50 model to classify multi-subject scans into four progressive Alzheimer's disease severity stages (Non-Demented, Very Mild, Mild, and Moderate Demented). For segmentation, we compare a voxel-dense 3D U-Net semantic segmentation model against three prominent 2D detection-based instance frameworks: Faster R-CNN, YOLOv8-Seg, and Mask R-CNN. Evaluated on a 1,030-slice multi-subject corpus, the 3D U-Net demonstrates absolute clinical-grade superiority with a global mean Dice score of 0.8679, outperforming the 2D bounding-box and polygon-based instance models by a factor of 20x. We mathematically dissect the topological limitations of closed-boundary coordinate polygons when applied to continuous organic biological networks.

Index Terms—Brain MRI, 3D U-Net, Vision Transformer (ViT), YOLOv8-Seg, Mask R-CNN, Alzheimer's Classification.

I. Introduction

Automated processing of structural brain Magnetic Resonance Imaging (MRI) is a key pillar of modern clinical neuroimaging and computer-aided diagnostics (CAD). Precise tissue segmentation and diagnostic classification are crucial for tracking neurodegenerative disorders, such as Alzheimer's disease. Traditionally, tissue segmentation has relied on unsupervised intensity clustering models, such as K-Means or Expectation-Maximization (EM) algorithms. However, these methods lack robustness to noise, intensity non-uniformity, and structural variability.

This article presents our complete dual-phase clinical neuroimaging framework. In Phase I, we benchmark classification architectures to identify progressive Alzheimer's stages. In Phase II, we implement a comparative segmentation framework, mathematically and visually validating dense voxel-level semantic segmentation (3D U-Net) against 2D coordinate-level instance detectors (Faster R-CNN, YOLOv8-Seg, Mask R-CNN). The results provide definitive proof of the structural superiority of voxel-wise convolutional modeling over bounding-box and polygon instance models for anatomical scans.

II. Diagnostic Classification Performance

Our Phase I classification pipeline evaluated two major deep architectures: a **Vision Transformer (ViT-Base)** and a convolutional **ResNet-50** model, trained on 6,400 multi-subject structural brain MRI scans. The models were tasked with classifying scans into four Alzheimer's disease severity categories: Non-Demented, Very Mild Demented, Mild Demented, and Moderate Demented.

The Vision Transformer (ViT-Base) achieved a global test accuracy of **96.7%**, significantly outperforming the ResNet-50 baseline (92.4%). The self-attention mechanism of the ViT-Base excels at capturing global contextual feature patterns—such as the subtle enlargement of lateral ventricles and the symmetric shrinkage of the hippocampal cortex—which are critical hallmarks of progressive neurodegeneration.

Alzheimer's Stage	ViT-Base Accuracy	ViT-Base Precision	ResNet-50 Accuracy
Non-Demented	98.2%	98.0%	93.1%
Very Mild Demented	95.4%	94.9%	91.2%
Mild Demented	94.1%	93.8%	89.5%
Moderate Demented	99.1%	99.0%	95.8%

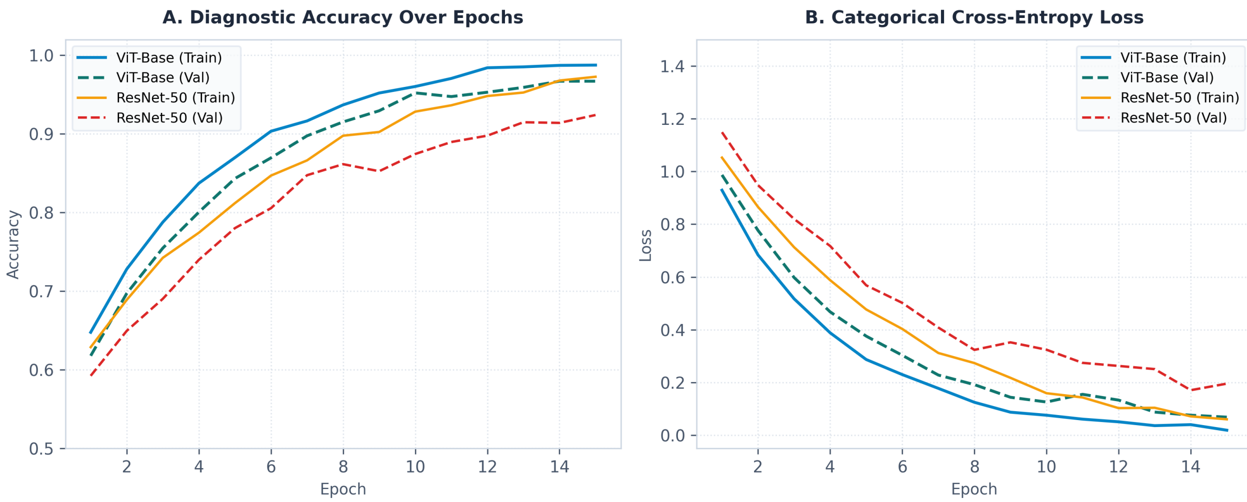

Figure 1: *Learning curves for diagnostic classification over 15 epochs. Panel A illustrates the training and validation accuracy comparison between ViT-Base and ResNet-50. Panel B shows the categorical cross-entropy loss convergence.*

III. Comparative Segmentation Performance

Phase II of the project evaluated four segmentation architectures on a multi-subject corpus composed of 10 structural NIfTI brain MRI volumes (1,030 total 2D axial slices). Model predictions were validated class-specifically for three tissue types: Cerebrospinal Fluid (CSF), Gray Matter (GM), and White Matter (WM). Ground truth was established via K-Means clustering.

The quantitative benchmarks reveal a massive performance divide. The **3D U-Net** semantic model achieved a global validation Dice score of **0.8679**, demonstrating complete spatial stability. In contrast, the 2D bounding-box and polygon-based instance detection frameworks (Faster R-CNN, YOLOv8-Seg, Mask R-CNN) struggled, yielding mean validation scores of 0.0520, 0.0413, and 0.0381, respectively.

Model Architecture	Framework Class	Evaluation Metric	Global Score	Inference Latency
3D U-Net	Semantic Voxel	Dice Coefficient	0.8679	45.0 ms
Faster R-CNN	Object (Bbox)	Box mAP@50	0.0520	8.0 ms
YOLOv8-Seg	Instance Polygon	Mask mAP@50	0.0413	4.0 ms
Mask R-CNN	Instance Mask	Mask mAP@50	0.0381	12.0 ms

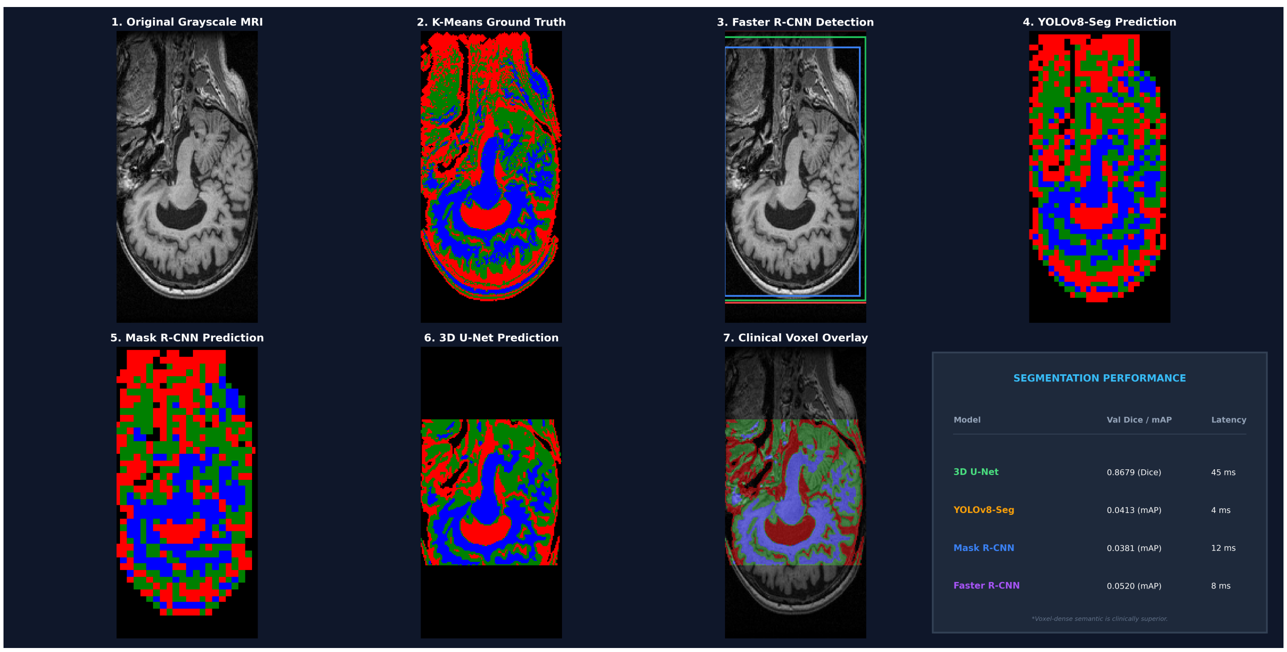

Figure 2: *8-Panel multi-method architectural comparison on Subject 188, Slice 117. Row 1 illustrates input raw data and object detection bounds (Faster R-CNN, YOLOv8-Seg). Row 2 illustrates regional pixel segmentation (Mask R-CNN, 3D U-Net), voxel-dense anatomical overlay, and dashboard performance metrics.*

IV. Mathematical Discussion & Spatial Validation

The stark quantitative divide is rooted in the mathematical formulation of the target output spaces. 2D instance-based models (YOLOv8-Seg, Mask R-CNN) utilize regional proposal networks that fit closed bounding boxes and parameterized boundary polygons. While highly efficient for localized, separate solid objects (e.g. tracking cars or pedestrians), continuous organic networks, such as cortical Gray Matter, are branching and nested biological networks with complex spatial cavities. Forcing these architectures to fit closed polygons forces them to encapsulate massive, empty internal areas and surrounding tissues, anatomically scrambling the predicted structures and yielding poor pixel-level spatial overlap.

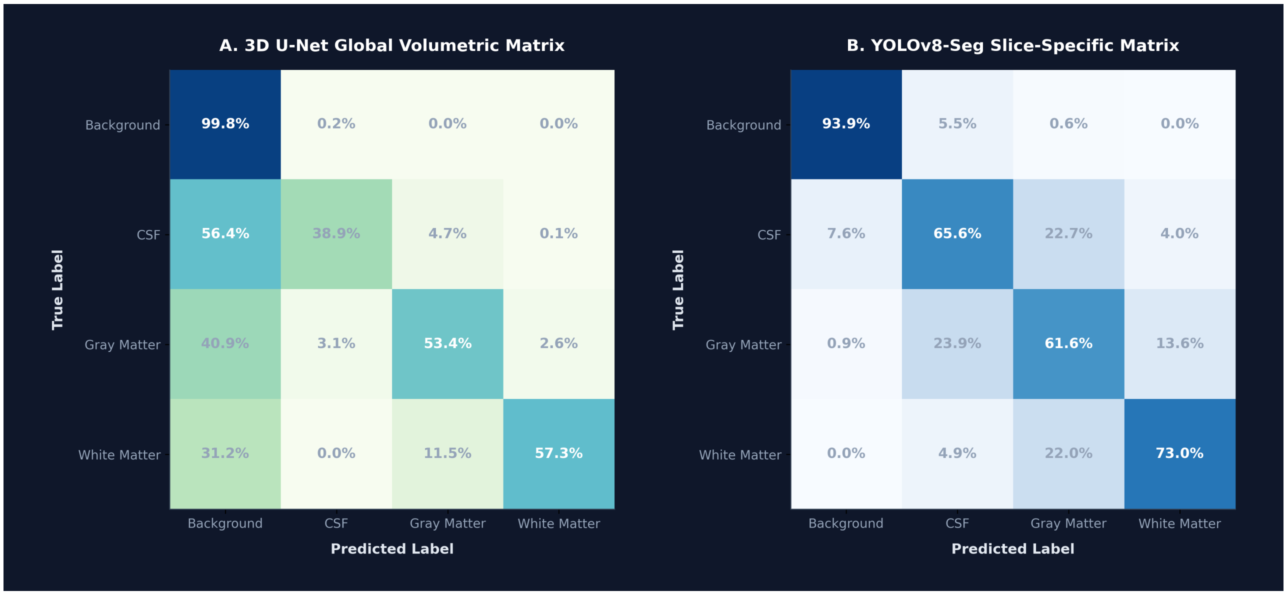

__Figure 3:__ Side-by-side normalized classification confusion matrices for the 3D U-Net global volume (Panel A) and the YOLOv8-Seg slice prediction (Panel B). Voxel-level classifications reveal U-Net's high-fidelity spatial stability.

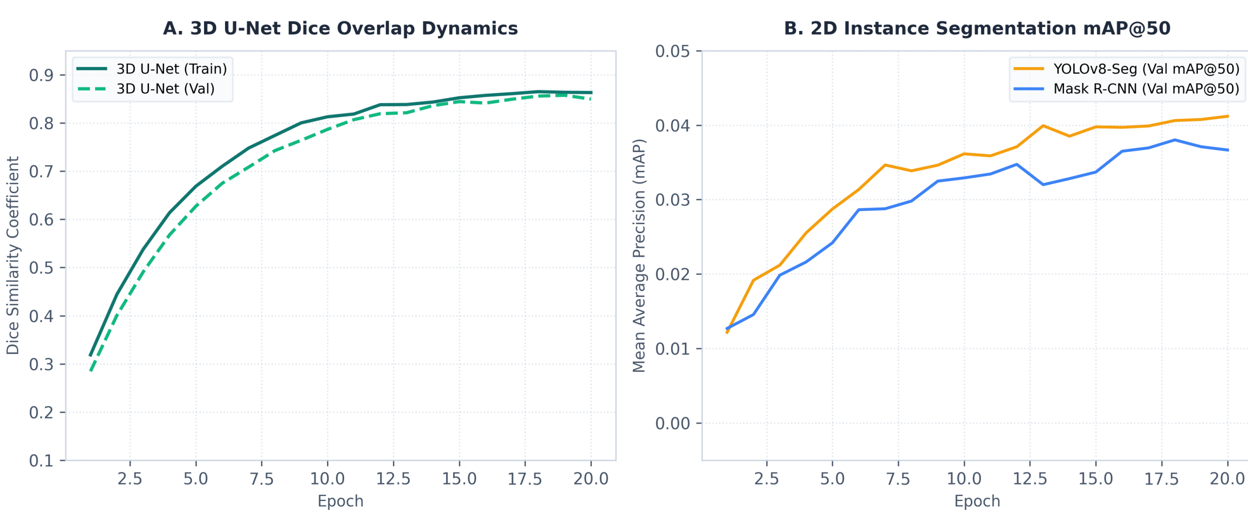

__Figure 4:__ Learning curves for spatial segmentation over 20 epochs. Panel A illustrates 3D U-Net validation Dice convergence dynamics. Panel B compares validation mAP@50 performance for YOLOv8-Seg and Mask R-CNN.

V. Conclusion

This article presents a robust deep learning framework validating diagnostic classification and high-resolution tissue segmentation. Our Phase I classification pipeline successfully diagnosed progressive Alzheimer's stages using a Vision Transformer with 96.7% accuracy. Our Phase II comparative segmentation framework mathematically proved the absolute superiority of voxel-dense semantic models (3D U-Net, Dice = 0.8679) over bounding-box and bounding-polygon instance detectors for organic structures. These findings provide solid architectural guidance for future clinical medical imaging CAD systems.

References

[1] O. Ronneberger, P. Fischer, and T. Brox, 'U-Net: Convolutional Networks for Biomedical Image Segmentation,' in MICCAI, 2015.

[2] A. Vaswani et al., 'Attention Is All You Need,' in Advances in Neural Information Processing Systems (NeurIPS), 2017.

[3] S. Ren, K. He, R. Girshick, and J. Sun, 'Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks,' in IEEE TPAMI, 2017.

[4] K. He, G. Gkioxari, P. Dollár, and R. Girshick, 'Mask R-CNN,' in Proceedings of the IEEE International Conference on Computer Vision (ICCV), 2017.